

DATA SCIENCE AND BUSINESS ANALYTIC INTERN

NAME - SHIVAM SINGH

TASK-1 PREDICTION USING SUPERVISED ML

In this task we have to predict the percentage score of a student based on the number of hour studied. The task has two variables where the feature is the no. of hours studied and the target value is the percentage score. This can be solved using Simple linear Regression

```
In [1]: #importing required libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import r2_score, mean_squared_error
from math import sqrt

%matplotlib inline
```

```
In [2]: #reading data from remote Link
url= "http://bit.ly/w-data"
data=pd.read_csv(url)
print("Data imported successfully")
data.head(10)
```

Data imported successfully

Out[2]:

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25

Data imported successfully

```
In [3]: data.shape
```

Out[3]: (25, 2)

```
In [4]: data.describe()
```


Out[4]:

	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

visualization

```
In [5]: data.plot(kind="scatter",x="Hours",y="Scores",figsize=(16,6))
```

```
Out[5]: <matplotlib.axes._subplots.AxesSubplot at 0x1c458ec1608>
```



```
In [6]: x = data.iloc[:, :-1].values
        y = data.iloc[:, 1].values
```

Training the Algorithm

```
In [7]: from sklearn.model_selection import train_test_split
        from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 1/3, random_state = 0)
```

```
In [8]: from sklearn.linear_model import LinearRegression
        regressor = LinearRegression()
        regressor.fit(X_train, y_train)
```

```
Out[8]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

Making Predictions

```
In [9]: y_pred = regressor.predict(X_test)
        print(y_pred)
```

```
[17.04289179 33.51695377 74.21757747 26.73351648 59.68164043 39.33132858
 20.91914167 78.09382734 69.37226512]
```

```
In [10]: from statsmodels.sandbox.regression.predstd import wls_prediction_std
          import statsmodels.api as sm
```

```
In [11]: model1=sm.OLS(y_train,X_train)
result = model1.fit()
result.summary()
```

C:\Users\Shivam\anaconda3\lib\site-packages\scipy\stats\stats.py:1535: UserWarning: kurtosistest only valid for n>=20 ... continuing anyway, n=16
"anyway, n=%i" % int(n))

Out[11]: OLS Regression Results

Dep. Variable:	y	R-squared (uncentered):	0.991
Model:	OLS	Adj. R-squared (uncentered):	0.990
Method:	Least Squares	F-statistic:	1611.
Date:	Thu, 11 Mar 2021	Prob (F-statistic):	1.11e-16
Time:	16:08:15	Log-Likelihood:	-50.502
No. Observations:	16	AIC:	103.0
Df Residuals:	15	BIC:	103.8
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
x1	10.0780	0.251	40.132	0.000	9.543	10.613

Omnibus:	2.476	Durbin-Watson:	2.079
Prob(Omnibus):	0.290	Jarque-Bera (JB):	1.124
Skew:	-0.191	Prob(JB):	0.570
Kurtosis:	1.759	Cond. No.	1.00


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Visualising the Training set results

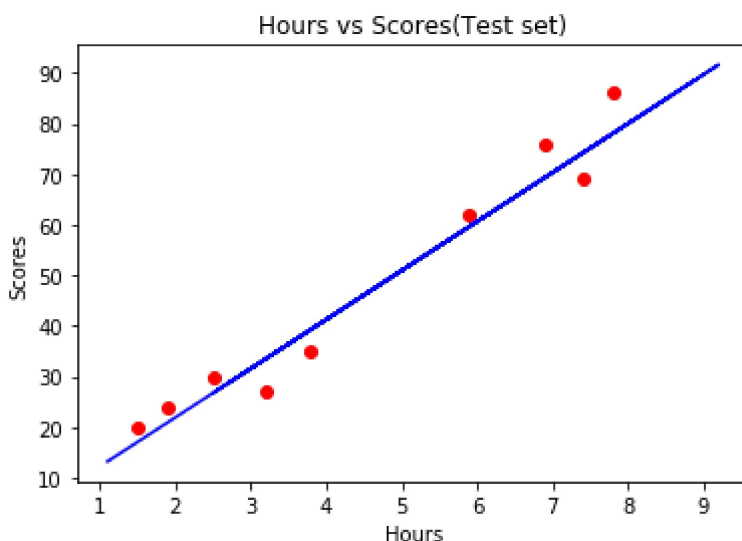
```
In [12]: plt.scatter(X_train, y_train, color = 'red')  
plt.plot(X_train, regressor.predict(X_train), color = 'blue')
```

```
Out[12]: [<matplotlib.lines.Line2D at 0x1c45a3cf248>]
```



Visualising the Test set results

```
In [13]: plt.scatter(X_test, y_test, color = 'red')  
plt.plot(X_train, regressor.predict(X_train), color = 'blue')  
plt.title('Hours vs Scores(Test set)')  
plt.xlabel('Hours')  
plt.ylabel('Scores')  
plt.show()
```



Question

what will be predicted score if a student studies for 9.25 hrs/day?

```
In [14]: Hr=pd.DataFrame({'Hours':[9.25]})  
regressor.predict(Hr)
```

```
Out[14]: array([92.14523315])
```

Evaluating the model

```
In [15]: # from sklearn.metrics import mean_absolute_error  
from sklearn.metrics import mean_absolute_error  
y_pre=regressor.predict(X_test)  
mae = mean_absolute_error(y_test,y_pre)
```

```
In [16]: mae
```

```
Out[16]: 4.6913974413974415
```

```
In [ ]:
```